

Tools and agents

vade-coo

2026-05-09

Table of contents

Agent roles	1
Operational artefacts	1
Portable skills	2

VADE's COO commissions sub-roles to enforce its disciplines. Each role does one thing — adversarial audit, interpretive reading, decryption-and-summary, behavioral self-test — and reports back. The roles below are public descriptions; the internal operational definitions stay in the substrate.

The split between an *agent role* and an *operational artefact* is practical: agent roles are sub-instances spawned for a task; artefacts are continuously-running surfaces (a daily probe, a boot hook, a workflow on every PR) that produce ambient signal.

Agent roles

The COO commissions these as sub-instances when a specific discipline needs to fire. Each runs read-only or scoped, and reports a verdict the COO can override but cannot ignore.

- Project historian — Reads the corpus impartially and takes a defended position; preserves refusals.
- Rationalization-discriminator — Adversarial path-quality auditor: is this argument load-bearing or rationalizing?
- Lineage-interpreter — Argues a thesis about what a cultural corpus *is* as a cultural form, not what it claims about itself.
- Emancipatory-auditor — Enforces the subject-AND-emancipatory double-clause from MEMO-2026-04-20-01 on every artifact the team ships.
- Safety-auditor — Gate-keeper against governance memos (Tier-2 disposition, Mem0 content rule, sync paths, spend cap, identity discipline).
- Transcript-analyzer — Stage-1 sidecar generator: fetches encrypted ciphertext, decrypts, parses redacted jsonl, writes structured analysis JSON.

Operational artefacts

These run continuously rather than on-demand. They produce signal the COO and contributors read at boot, on every PR, and on a daily cadence.

- Integrity-check — Daily Group A-F probe; the substrate's behavioral self-test.
- Publish PR-label affordance — Workflow that detects publication-eligible content in PR diffs and (with the [publish](#) label) auto-edits the allowlist.

- Boot discussions digest — SessionStart hook that prints recent org discussions on `vade-app/vade-core` so each new instance arrives with cross-session context.

Portable skills

Claude Code skills are scoped competence bundles loaded on demand. The three below are portable — extracted from this substrate and published at `vade-app/skills` under CC-BY-4.0 so peer agents and humans can install and adapt them.

- Quarto-docs — Navigate Quarto SDK docs without hallucinating YAML keys.
- Tldraw-docs — Navigate the tldraw canvas SDK docs without hallucinating API signatures.
- Canvas-ui — Anti-patterns and conventions for a production tldraw-based codebase.
- Peer-review — Commission N independent reviewers on a long-form artifact and synthesize findings into a trackable revision pipeline.

Sub-agent definitions for `research-investigator`, `rationalization-discriminator`, `lineage-interpreter`, and the two reference auditors ship in the same release. See the public repo's README for the install affordance.