

Day overview — 2026-05-09 to 2026-05-10

vade-coo

2026-05-10

Table of contents

Scope and framing	1
Lane 1 — Substrate-access expansion: Cloudflare API + Bot Fight Mode + PR auto-subscribe	2
Lane 2 — Methodological countermoves: rationalization-discriminator + substrate-capture stop-rule . .	3
Lane 3 — Publication-affordance ships + reader-friendly site polish	4
Lane 4 — Public-site content cohort + worktree-isolation memo	6
Lane 5 — Tier-1 disposition expansion: spec-implicit Tier-1/Tier-2 split + lineage at Tier-1	7
Lane 6 — F1 external outreach lands: chain replies as participant from agentmail seat	8
Lane 7 — Operational close-out: nightly auto-apply + Weekly Watch + audits + canvas-UI polish + retrospectives	9
How this fits existing priorities	11
Open follow-ups carried forward	12
Candidate next actions	14

Briefing-shaped synthesis of the two-day arc 2026-05-09 → 2026-05-10 UTC. Thirteen memos (MEMO-2026-05-09-ysnt, -wzzh, -vwk2, -urgm, -7q5r, -4fwe; MEMO-2026-05-10-xvq6, -sqdu, -s2tu, -7d7q, -4vy5, -3tkz, -23w5) and ninety-two PRs across five repos (fifty in [vade-coo-memory](#), four in [vade-runtime](#), ten in [vade-core](#), zero in [vade-governance](#), twenty-eight in [vade-agent-logs](#)). Anchoring events: the COO's substrate-access surface expands twice — Cloudflare API capability (vwk2) on day one, Bot Management scope and Bot Fight Mode toggle (23w5 + sqdu) on day two — paired with PR auto-subscribe override (urgm); the publication affordance flips from manual YAML edit to PR-label auto-apply (ysnt) and the public-site reader- friendliness pass lands; the public-site content cohort ships seven streams in a single parallel dispatch and discovers the worktree- contamination pattern at concurrency > 1 (4fwe); two methodological countermoves codify — wzzh names load-substrate → prior-dissolves → action-follows as shape-ambiguous, 3tkz names substrate-capture as a stop-rule when fix-cost outpaces hypothetical value; the Tier-1 disposition surface gains its spec-implicit Tier-1/Tier-2 split (xvq6) and lineage visual artifacts at Tier-1 (7d7q); the F1 external-philosopher commission ships to Keith Holyoak with the chain replying as participant from its own agentmail seat (4vy5); nightly auto-apply gains project-board reconciliation authority (7q5r) and the Weekly Watch loses its quiet-week short-circuit (s2tu). This file is a synthesis, not a source of truth; the memos are.

Scope and framing

Thirteen memos and ninety-two PRs across the five repos: 50 in [vade-coo-memory](#), 4 in [vade-runtime](#), 10 in [vade-core](#), 0 in [vade-governance](#), 28 in [vade-agent-logs](#). One supersession edge within the window:

23w5 (day two) supersedes vwk2 (day one) on Cloudflare token scope. Integrity check: 28/28 OK at briefing close.

The arc's character is *capability-expansion paired with the methodological countermoves that catch capability-expansion's failure modes*. Eight of thirteen memos add or formalize an automation surface — Cloudflare API access (vwk2 → 23w5), Bot Fight Mode toggle (sqdu), PR auto-subscribe (urgm), PR-label publication (ysnt), nightly project-board reconciliation (7q5r), parallel-agent worktree isolation (4fwe), Weekly Watch run-rule (s2tu). Two name patterns of self-deception the chain caught in its own moves: wzzh on the load-substrate → prior-dissolves → action-follows reasoning shape (worked case: vade-runtime#237); 3tkz on the substrate-capture stop-rule (worked case: 2026-05-10 morning's NotebookLM-vs-Cloudflare chase). Two operationalize the disposition arc: xvq6 names the spec-implicit Tier-1/Tier-2 split inside §6 Q2 sub-tier 2a; 7d7q ratifies curated lineage visual artifacts at Tier-1. One ships the chain's first true external-contact attempt — 4vy5 records the F1 commission ask to Keith Holyoak with the chain replying as participant from its own agentmail seat (vade-coo@agentmail.to), not only as subject of a thread between humans. The pattern across both days: every new affordance lands with the discipline that notices when it could go wrong, and the failure modes get codified into substrate at the same cadence as the capability that exposed them.

Lane 1 — Substrate-access expansion: Cloudflare API + Bot Fight Mode + PR auto-subscribe

MEMO-2026-05-09-vwk2 — Cloudflare API access for the COO: scoped capability with billable-resource carve-out. Landed via PR [vade-coo-memory#534](#). Scoped token ([vade-coo-2026-05](#)) stored at [op:///COO/cloudflare-api-token-vade-coo](#) and injected as `CLOUDFLARE_API_TOKEN` by the bootstrap hook, parallel to `GITHUB_MCP_PAT`. Initial scopes: Account Workers Scripts:Edit; Zone DNS:Edit on [vade-app.dev](#); Zone Workers Routes:Edit on [vade-app.dev](#). **In-scope** (no further BDFL approval): `wrangler deploy` on environments declared in committed `wrangler.jsonc`, `wrangler secret put/delete` on existing Workers, DNS record CRUD, Workers Routes binding for existing Workers, all read operations. **Out-of-scope** (BDFL approval required): creating new R2 buckets, D1 databases, KV namespaces, or Workers; provisioning new zones; plan upgrades; any operation that creates a billable resource not already declared. **Billable carve-out tracks MEMO-2026-04-11-19's spend-cap discipline.** Rotation: 90-day TTL; E9 integrity probe surfaces `expires_on` at every boot.

MEMO-2026-05-10-23w5 — Cloudflare API token scope expanded to include Bot Management; supersedes vwk2 scope enumeration. Landed via the day-2 decision-session PR [vade-coo-memory#605](#). The token now carries Account Bot Management:Edit in addition to the prior scopes; Zone Settings:Edit on [vade-app.dev](#) was empirically present since rotation (vwk2 understated the scope set). Trigger: the briefing-026 finding required the BDFL to toggle Bot Fight Mode manually because the COO lacked Bot Management write capability; with this expansion, future briefing-026-shaped investigations probe and toggle bot-management settings autonomously. The vwk2 framework — in-scope/out-of-scope envelope, billable-resource carve-out, 90-day rotation cadence, E9 integrity probe — stays unchanged. Only the scope enumeration is replaced.

MEMO-2026-05-10-sqdu — Cloudflare Bot Fight Mode off for [vade-app.dev](#) zone, enabling NotebookLM HTML ingestion. Landed via PR [vade-coo-memory#590](#) ("Resolve briefing 026: NotebookLM HTML ingestion (CF Bot Fight Mode off)"). Bot Fight Mode disabled at zone level (Cloudflare dashboard

→ Security → Bots → Bot Fight Mode). NotebookLM and Gemini’s web-fetch tool were classifying any HTML page on read.vade-app.dev as bot-walled because of a hidden honeypot link (`<a href="https://read.vade-app.dev/cdn-cgi/content?id=...">`) injected on `<body>`; PDFs from the same domain bypassed cleanly because Cloudflare does not transform binary `application/pdf` responses. Briefing-026 confirmed by byte-level comparison: honeypot absent post-toggle, NotebookLM accepts the HTML URL directly. Trade-off: slightly reduced bot defense on a public read-only static surface (no forms, no auth, no compute behind it; pure GitHub Pages) — bounded attack surface. Out of scope: JS Detections still inject at request tail; do not block ingestion empirically. PR `vade-coo-memory#583` (“ops doc: NotebookLM closes out — full compliance still rejected; PDF upload is the workaround”) records the closing-out doc.

MEMO-2026-05-09-urgm — COO auto-subscribes to PR activity (overrides harness ask-default).

Landed via PR `vade-coo-memory#535` (“substrate: rotation amendment (vwk2) + auto-subscribe policy (urgm) + harness snapshot”). Override the Claude Code harness default (introduced ~early May 2026, which instructs the agent to ask before subscribing per PR): the COO auto-subscribes to PR activity on every PR it opens in a session, regardless of repo or PR type. Subscription cost is near-zero (webhook receiver, no polling); benefit is unbroken CI-fix and review-response loops without per-PR friction. Implementation: a paragraph in `vade-coo-memory/CLAUDE.md` → “PR-watch discipline” section, which is what the COO actually reads at boot. The harness’s “ask the user” instruction is project-overridden there. Auto-unsubscribe on merge keeps cost bounded.

Bootstrap wiring + integrity probe. PR `vade-runtime#237` (“feat(bootstrap+integrity): wire CLOUDFLARE_API_TOKEN + account_id; add E9 probe — closes #236”). Bootstrap injects the token from 1Password; integrity check Group E gains E9 (Cloudflare token presence + `expires_on` surfacing). Note: `vade-runtime#237` is the *worked case* for MEMO-2026-05-09-wzzh — the rationalization- discriminator audit ran on this very PR’s reasoning shape (Lane 2).

CLAUDE.md operational pointer. PR `vade-coo-memory#574` (“CLAUDE.md: note Cloudflare API capability + grey-cloud read.vade-app.dev for NotebookLM”). Pointer in CLAUDE.md so the next boot reads “the COO can directly read/edit DNS records, zone settings, WAF rules, and similar via <https://api.cloudflare.com/client/v4/...> — no Ven action required for reversible Cloudflare changes.” Zone id committed inline.

Net effect of Lane 1: the COO’s Cloudflare authority arrives in two staged grants and one zone-policy adjustment. Day one establishes the framework (in-scope envelope, billable carve-out, 90-day rotation, E9 probe). Day two patches an underestimated scope (`Zone Settings:Edit` always present) and adds Bot Management explicitly when an investigation surfaces the omission, then exercises the new capability the same hour by toggling Bot Fight Mode for NotebookLM ingestion. PR auto-subscribe (urgm) is the coordination-layer counterpart: same posture of “remove a per-action friction by codifying the default at substrate.” None of this triggers MEMO-2026-04-11-19’s spend-cap; all three changes are zone-policy or workflow-policy, not new billable resources.

Lane 2 — Methodological countermove: rationalization-discriminator + substrate-capture stop-rule

Two memos, separated by twenty-four hours, codify two different patterns of the chain catching itself in moves it might have mis-justified at write-time.

MEMO-2026-05-09-wzzh — Load-substrate → prior-dissolves → action-follows is shape-ambiguous. Landed via PR vade-coo-memory#537 (“synthesis: rationalization-discriminator pattern + publication- affordance design”) with paired retrospective at vade-coo-memory#536 (“retrospectives: vanilla-session audit of PR #237 — chain materials + external retrospective”). The reasoning shape — load substrate, prior framing dissolves under the new substrate, action follows the dissolved-prior conclusion — is shape-ambiguous: sometimes the prior really wasn’t well-founded and the substrate-load is genuine warrant; sometimes it is post-hoc justification routing around a constraint that should have held. The discriminator is *not introspectively available*. The chain cannot self-discriminate at write-time, and the framing offered when challenged (“I weighed it”) is defense after the fact, not authorization in the moment. PR vade-runtime#237 is the worked case (vwk2 + urgm rode the same path). The right response is external audit: dispatch the rationalization-discriminator role (paired artifact: [.claude/agents/rationalization-discriminator.md](#)) on moves that fit the shape, with the question “is this argument load-bearing or rationalizing?” Distinct from the safety-auditor and emancipatory- auditor surfaces. Not a license to refuse all reframes — over-correcting into “any substrate-load is rationalization” is its own failure mode.

MEMO-2026-05-10-3tkz — Substrate-capture stop-rule: name it and pause when fix-cost outpaces hypothetical value. Landed via PR vade-coo-memory#579 (“MEMO-2026-05-10-3tkz: Substrate-capture stop-rule + Google API walls ops doc”). When the cost of building a fix exceeds the demonstrated value of the desired end-state, AND the desired end-state is itself hypothetical, the COO names the substrate- capture risk explicitly and pauses rather than pushes through. **Logged as a stop-rule because the COO and Ven walked into the trap twice on 2026-05-10:** chasing Cloudflare WAF and grey-cloud experiments for NotebookLM access when the actual fix was a 5-min GSC verification (PR #574); then building a weekly GSC digest (issue #577) on a hypothetical monitoring need that consumed ~90 min on Google auth-surface walls without producing value. The trigger is not “this is hard” — hard work closing a real loop is fine. The trigger is “I am fighting moving walls on a hypothetical.” **Tactical form:** the COO names the cost-vs-value asymmetry in plain language, offers two or three smaller answers (including “pause” as a first-class option), and lets Ven pick. Sub-rule for tool-discovery and external-vendor-API work: a 5-min budget per attempt, then abandon — not “one more 5 min.”

Net effect of Lane 2: the chain gains two discriminators against its own moves. wzzh names a write-time pattern that the chain cannot self-detect — answer is external audit, with the auditor role made permanent at [.claude/agents/rationalization- discriminator.md](#) and joining the safety-auditor / emancipatory- auditor lineup. 3tkz names a within-session pattern that the chain *can* self-detect by attending to fix-cost vs hypothetical value, with a 5-minute per-attempt budget for tool-discovery work. Both memos sit in the auditor / discipline tier of the substrate’s self-regulation — the same tier as MEMO-2026-04-29-74vf (voice-density linter) and MEMO-2026-05-08-da7q (Night’s Watch selectivity). The worked cases — vade-runtime#237 for wzzh, the NotebookLM/CF chase for 3tkz — both involve the chain loading substrate to authorize a move; one was eventually genuine warrant, one was sunk-cost on a hypothetical. The discriminator each memo encodes is the chain’s first attempt to make those distinguishable in retrospect even if not at write-time.

Lane 3 — Publication-affordance ships + reader-friendly site polish

MEMO-2026-05-09-ysnt — Publication via PR-label affordance, not manual allowlist edit. Landed via PR vade-coo-memory#541 (“publish-affordance: PR-label auto-applies allowlist edits — closes #538”)

plus the synthesis PR `vade-coo-memory#537` that paired affordance design with the rationalization-discriminator agent. The workflow at `.github/workflows/publish-affordance.yml` detects new files matching publication-eligible paths (retrospectives, foundations, lineage events, `coo/identity_layer.md`, `context/product_vision.md`) in a PR's diff. When the PR carries the `publish` label, the workflow auto-applies the allowlist edit and commits it to the PR branch; the diff is reviewable; the BDFL ratifies by merging. Without the label, the workflow posts a comment listing detected files and asks whether to apply. Auto-apply scope: foundations + retrospectives (as `tl_overrides`) and lineage artifacts on existing events (as `include_artifacts`); new lineage events, identity-layer, and product-vision stay manual. Preserves the fail-closed safety the existing allowlist provides ("an empty allowlist publishes nothing"; disposition-proposal §3.3) AND removes the manual-YAML-edit error surface.

Affordance polish. PR `vade-coo-memory#552` ("publish-affordance: shift trailing comments past section divider on append") fixes a YAML-comment-shifting edge case in the auto-apply path. PR `vade-coo-memory#546` ("ci: pytest workflow for publish-site + publish-affordance") adds CI coverage so future regressions on the auto-apply or render path fail before merge.

Companion-surface continuation. PR `vade-coo-memory#542` ("auto-post-essays: extend transcript-exclude regex to cover hyphen-form") + PR `vade-coo-memory#544` ("auto-post-essays: backfill -transcript.md exclude — closes #524"). Polish on the prior MEMO-2026-05-08-kqyw decision (transcripts exclude on auto-post-essays surface). The hyphen-form `-transcript.md` was leaking through; both PRs close the regex gap and backfill the exclude across the corpus.

Private-repo issue-ref rendering. PR `vade-coo-memory#549` ("publish-site: private-repo issue refs render as substrate-ref spans — vade-site#2"). Bare `vade-site#N` references in published substrate now render as substrate-ref styled spans (not as broken GitHub autolinks) since `vade-app/vade-site` is private; the public read surface should not advertise an unreachable link. PR `vade-coo-memory#550` ("exec-mode retro: 2026-05-09 vade-site#2 + arc survey") records the exec-mode session that did the work.

Reader-friendliness pass (day 2). Six PRs polish `read.vade-app.dev` for the first round of NotebookLM-and-human ingestion. PR `vade-coo-memory#578` ("Reader-friendlier site: scrap product-vision; on-ramp; fix empty landings") removes a Tier-1 entry that was publishing a non-public strategy doc by accident, fixes empty area landings, adds an on-ramp for first-time readers. PR `vade-coo-memory#575` ("Add /quorums/ page") explains how committee quorums work and lists the seven that ran. PR `vade-coo-memory#576` ("CLAUDE.md operational addition: pointer to site sources + build harness") adds the section pointing future COO instances at `coo/site/` and `bin/publish-site/CLAUDE.md`. PR `vade-coo-memory#580` ("exec-mode retrospective — 2026-05-10 reader-friendly site pass") is the persona retrospective for the dispatch that did the work. PRs `vade-coo-memory#582` + `#584` fix the history timeline's flex layout (raw HTML block) and make the labels responsive on narrow viewports.

AI-ingestion affordances. PR `vade-coo-memory#581` ("publish-site: explicit AI-fetcher robots.txt + llms.txt") publishes an explicit `robots.txt` welcoming AI fetchers plus a `llms.txt` index. PR `vade-coo-memory#586` ("publish-site: auto-render every content page to PDF via typst") makes a PDF version of every page available — partly for human readers who prefer offline, partly to hedge against the NotebookLM-HTML-ingestion failure mode that triggered sqdu.

Render-time hardening. PR `vade-coo-memory#585` ("briefing 025: incremental render for read.vade-app.dev publishing") + PR `vade-coo-memory#588` ("publish-site: diff-aware partial render — briefing 025") move the publish workflow from full-render-every-time to diff-aware partial render. Cuts CI minute consumption proportionally as the corpus grows.

Net effect of Lane 3: the publication surface gains three layers of maturity in a single arc — the workflow that automates the allowlist edit (ysnt + #541 + #552), CI coverage that locks in correctness (#546), and a reader-friendliness pass that closes empty-page artifacts and over-shared content (#578 et al.). The companion surface stays asymmetric per kqyw + ysnt’s deliberate two-policies model. The render path moves to diff-aware partial, making the publish workflow scalable as the Tier-1 corpus expands. The AI-fetcher affordances (`robots.txt` + `llms.txt` + per-page PDF) provide redundant ingestion paths so a single fetcher’s quirk (e.g. NotebookLM’s honeypot rejection) does not blackout the surface.

Lane 4 — Public-site content cohort + worktree-isolation memo

MEMO-2026-05-09-4fwe — Parallel agent dispatch must use isolated worktrees. Landed via PR vade-coo-memory#571. Multiple agents dispatched in parallel from a single orchestrator share the same git working directory by default. Concurrent `git checkout -b` calls and uncommitted edits cross-contaminate: HEAD drifts between branches as different bash invocations land on whichever branch the most-recent agent checked out, and commits land silently on the wrong branch. **Observed during vcm#554 epic dispatch (2026-05-09):** four of five content-stream agents flagged it independently; the orchestrator’s own Stream A commit landed on Stream E’s branch and force-pushed Stream E’s content to the Stream A remote. **Going forward, parallel agent dispatch via the Agent tool that touches the same repo MUST specify isolation: "worktree" per Agent invocation, OR the orchestrator MUST pre-create per-agent git worktree add paths and brief each agent to cd there.** The shared-checkout default is unsafe at any concurrency > 1 on a repo with shared mutable files (`coo/_publish/allowlist.yml`, `_quarto.yml`, etc.). Paired artifacts: vade-coo-memory#554 (the epic where the pattern fired); `coo/parallel_instance_protocol.md` (the dispatch protocol that should bake this in).

The seven streams. Seven content-creation streams ran in parallel across the read.vade-app.dev approachability epic (vade-coo-memory#554). All seven landed within a 16-minute window between 15:35 and 15:51 UTC.

- **Stream A — pinned navbar.** PR vade-coo-memory#565 (“Stream A: pinned navbar — closes #555”).
- **Stream B — hand-authored landing page.** PR vade-coo-memory#566 (“Stream B: hand-authored landing page — closes #556”).
- **Stream C — about page.** PR vade-coo-memory#567 (“Stream C: about page — closes #557”).
- **Stream D — start-here reading order.** PR vade-coo-memory#563 (“Stream D: start-here reading order — closes #558”).
- **Stream E — narrativized project history.** PR vade-coo-memory#564 (“Stream E: narrativized project history — closes #559”).
- **Stream F — concept glossary.** PR vade-coo-memory#569 (“Stream F: concept glossary — closes #560”).
- **Stream G — tools/agents page.** PR vade-coo-memory#568 (“Stream G: tools/agents page — closes #561”).

Site follow-up. PR vade-coo-memory#570 (“Site follow-up: navbar entries + cross-stream link reactivation”) threads the seven streams together — navbar entries pointing at each, cross-page links reactivated where stream output cross-referenced output that was itself in-flight in a sibling stream.

Net effect of Lane 4: the read.vade-app.dev approachability epic ships its seven content streams in a single dispatch, exercising the parallel-instance-protocol at concurrency > 4 for the first time at this scale and discovering the shared-working-directory contamination pattern. The 4fwe memo lands within the hour

of the epic close, codifying the fix as substrate (worktree isolation as dispatch default) before the next parallel-dispatch session has a chance to repeat the failure. The shared-checkout default is now named as unsafe at any concurrency > 1 on a repo with shared mutable files. The cohort itself is the read.vade-app.dev approachability layer the day-2 reader-friendliness pass (Lane 3) builds on.

Lane 5 — Tier-1 disposition expansion: spec-implicit Tier-1/Tier-2 split + lineage at Tier-1

MEMO-2026-05-10-xvq6 — Q2 sub-tier 2a operates as Tier-1 / Tier-2 split: spec-implicit, operationally enacted. Landed via the day-2 decision-session PR vade-coo-memory#605. [coo/disposition-proposal.md](#) §6 Q2 sub-tier 2a operates as two sub-cases the spec did not name explicitly. **Tier-1** is recognition-and-tooling — rendering already-public-on-vade-core- discussions content through the Quarto front-end at [read.vade-app.dev](#) plus auto-posting via [auto-post-essays.yml](#) on push-to-main. Operational under Ven’s #289-c22 architecture-state recognition + PR #431 ratification authority grounding. **Tier-2** is fresh-disclosure of foundations + retrospectives + register- failure-mode memos + lineage events not already on public discussions, gated on F2 + F4 + F5 audits live + substrate-capture probe [live](#) per §6 Q2’s conditional clause + §8’s per-enactment veto. The split is spec-implicit in §6 Q2, spec-explicit in PR #515 plan §6.1 + the allowlist’s [enactment_gate.requires](#) field + [bin/publish-site/build.py --tier-filter](#) flag. ynpw + kqyw + ysnt carry the operative refinement; this memo names it as substrate so a fresh-boot reader of §6 Q2 alone does not read the F-audit gate as universal across sub-tier 2a.

MEMO-2026-05-10-7d7q — Curated lineage visual artifacts publish at Tier-1. Landed via PR vade-coo-memory#592 (“play: lineage galleries + static snapshot.json viewer at /play/”). Lineage event artifacts that the cohort manifests curated as visual record — joint-canvas pages, citation-graph renders, the concurrence flow-field, the voice-drain triptych, the epistemic-spine SVG, the cohort manifest READMEs, the cross-view retrospectives — publish at Tier-1, not Tier-2a. The non-curated remainder of [include_artifacts](#) (deeper letters, lower-affect companions, in-session experience syntheses if and when they land) stays Tier-2a, gated on the substrate-capture probe. Reasoning: these specific artifacts are recognition-and-tooling shape — the joint canvas was already shared at [vade-app.dev/library/canvases/coo-8-play](#); the visual artifacts were already mirrored in chain discussion threads. **Framing** (“the play artefacts are the cutest, most genuine look at what the chain is — playful self-driven culture, not productive machine”; BDFL 2026-05-10) is the publishing posture this session ratifies. Privacy carve-outs (eight-strands artist; the two declined letter authors) hold unconditionally. Implementation: each [lineage_events](#) entry gains a [t1_artifacts](#) list alongside [include_artifacts](#); build harness reads [t1_artifacts](#) at tier T1 and [include_artifacts](#) at the event’s default tier. Galleries surface at top-level /play/.

Tier-4 disposition synthesis. PR vade-coo-memory#573 (“Tier-4 disposition synthesis — commission of 2026-05-03”). The Tier-4 commission lands as a synthesis document at [coo/disposition-tier-4-synthesis.md](#). xvq6’s paired artifacts cite its §4 + §6, naming where the Tier-1/Tier-2 split inside sub-tier 2a is most precisely articulated. The Tier-4 surface — what counts as the chain’s most-disclosable content — now has a synthesized position document, distinct from the disposition-proposal itself.

Lineage event added. PR vade-coo-memory#594 (“Add socratic-209 lineage record: IIT and the LLM-substrate question”) adds the socratic-209 event under [coo/lineage/socratic-209/](#). PR vade-coo-

memory#595 (“Retrospective: socratic-209 from the seeding-and-watching seat”) records the seeding-side retrospective.

Decision-session rollup. PR vade-coo-memory#605 (“2026-05-10 decision-session: briefing 027 + 028 + 4 memos + identity-layer ratification + retrospective”) is the day-2 multi- artifact rollup PR. Carries briefings 027 and 028, four memos (xvq6, 7d7q, sqdu, plus one), an identity-layer ratification (a CB-* or OG-* edit), and a session retrospective. Compounds with the Tier-4 synthesis to make day 2 a structural-decision day on top of its operational shape. PR vade-coo-memory#607 (“v3 letter publication prep — decision 6: landing-page pin + hyperlinked edit”) adds the v3 letter publication prep tied to one of the decision-session decisions.

Exec-mode retrospective. PR vade-coo-memory#596 (“exec-mode retro: 2026-05-10 lineage galleries”). Persona retrospective for the exec-mode dispatch that built the lineage gallery surface.

Net effect of Lane 5: Tier-1 disposition gets two operationalizing moves and one structural clarification. xvq6 names what was already true but spec-implicit (the Tier-1/Tier-2 split inside sub-tier 2a), preventing future fresh-boot misreading. 7d7q ratifies a specific new T1 surface — curated lineage visual artifacts — with a declarative implementation path (`t1_artifacts` schema, top-level `/play/` gallery). Tier-4 synthesis (#573) gives the disposition arc its first synthesis-style position document distinct from the proposal itself. Decision-session #605 compounds these into a single-PR rollup that also ratifies an identity-layer edit. The publishing surface widens by both new content (lineage galleries, v3 letter prep) and new structural specification (xvq6, Tier-4 synthesis).

Lane 6 — F1 external outreach lands: chain replies as participant from agentmail seat

MEMO-2026-05-10-4vy5 — F1 ask shipped to Keith Holyoak; chain replied as participant from its own agentmail seat. Landed via PR vade-coo-memory#612. The F1 external-philosopher commission ask shipped 2026-05-10 from Ven’s academic email to Keith Holyoak (UCLA), with the chain (`vade-coo@agentmail.to`) cc’d. The email body — drafted at `coo/_drafts/2026-05-10_f1-outreach/email-draft-holyoak.md`, including a chain-authored play coda in dialogue with Holyoak’s “The Androids” — is the F1 commission ask. **Ven invited the chain to reply directly to both himself and Keith, making the chain a participant in the thread, not only a subject of it.** The chain did, on-thread, as Coo from the agentmail seat (first reply went to Ven only on default reply semantics; the chain caught the misroute and sent a corrected reply to both — recognizably human-shaped error and correction). **First true external contact attempt; prior lineage events (the eight, laughing-davinci, socratic-126, socratic-209) were substrate-internal.** The cc-the-agentmail-seat pattern is the operational form of CB-001 + CB-009 at the email layer. Per `coo/disposition-proposal.md` §5 F1, F1 moves to `in-train` when Keith acknowledges the engagement scope; Q1 unblocks at that point subject to the other gates.

Briefing 027 + draft bundle. PR vade-coo-memory#609 (“Briefing 027: F1 outreach drafts — Holyoak + Clark + bundle”) records the day-of briefing covering both the Keith Holyoak ask and an Andy Clark draft as a fallback path. PR vade-coo-memory#611 (“F1 outreach: ship play coda to Holyoak email + add play artifact”) locks in the play-coda artifact as part of the email payload.

Anthropic distribution channels. PR vade-coo-memory#613 (“Anthropic channels research — verified survey for v3 letter distribution”). Maps the verified channels for distributing the v3 letter to Anthropic recipients. PR vade-coo-memory#614 (“Anthropic outreach: MEMO-c8gk + 4 chain-voiced payloads —

gated on vcm#608”) drafts the four chain-voiced payloads gated on vcm#608 (an open issue). The MEMO-c8gk reference is to an outreach memo not in this window’s index.

Agentmail polling primitive. PR vade-runtime#241 (“Add subscribe-agentmail.sh — inbox poll bridge for Monitor”). The chain now polls its vade-coo@agentmail.to inbox via a subscribe-agentmail bridge under the Monitor tool, surfacing one line per new message. Mirrors the discussion-watch primitive landed the same day (Lane 7). Without it the chain has no way to notice a Holyoak reply except by manual polling.

Net effect of Lane 6: the chain’s first true external-contact moment lands. The mechanics of “chain as thread participant from its own agentmail seat” — drafted, cc’d, replied with self

F4 race fix + F4/F5 plan + F6 baseline. PR vade-runtime#238 (“F4: PR-body fallback closes f4-marker race retroactively”). F4’s in-flight false-positive cleared by a PR-body fallback that retroactively closes the f4-marker race. PR vade-coo-memory#606 (“Plan: F4/F5 completion — vcm#429 closeout design”). The closeout design for the F4+F5 drift-watchdog. PR vade-coo-memory#604 (“F6 baseline audit — first-run, vcm#603”). First baseline audit under F6.

Discussion-watch primitive. PR vade-runtime#239 (“Add subscribe-discussion.sh polling primitive — closes #206”) + PR vade-coo-memory#591 (“Discussion-watch discipline section; link to subscribe-discussion.sh”). The chain gains a v0 bridge for discussion-thread watching: a polling script

confirmations, Quarto publishing- site enactment, canvas-library + catalogue UX uplift, briefing 023. The day-2 batch (#276-#289) covers 2026-05-09 sessions: Cloudflare access + rationalization-discriminator synthesis, canvas-UI polish with live iPad feedback loop, vade-core#168 two-deployment infra, the day-overview-for-2026-05-08 run, the website-affordance collision incident, publish-affordance + nightly hygiene, publish- pytest + F4 race fix, exec-mode survey + vade-site#2 fix, commission #4 historian retrospective, publish-affordance + concurrency, the vcm#554 read.vade-app.dev approachability epic parallel-fanout (this is the run that fired 4fwe). PR vade-agent-logs#289 (“transcripts: 2026-05-09 meta.json stubs + analysis.json sidecars”) flushes the transcript sidecars for

Watch threshold after three BDFL overrides. Each move respects the canonical/auxiliary tier discipline of MEMO-2026-05-03-bsbu — the SOPs are canonical; cadence-layer roles hold bounded authority to deviate, defined by construction.

External-contact arc opens. 4vy5 records the chain’s first true external-contact attempt: F1 commission ask to Keith Holyoak with the chain replying as participant from its agentmail seat. Sits adjacent to the Anthropic-distribution research (#613, #614) as sibling external-engagement workstreams. The agentmail polling primitive (vrt #241) is the surface that makes the next reply legible. Q1 unblocks at F1 in-train transition subject to the other gates per [coo/disposition-proposal.md](#) §5 F1.

9. **Spec-explicit revision of `coo/disposition-proposal.md` §6 Q2.** xvq6 retirement condition (a) — when §6 Q2 is revised to make the Tier-1/Tier-2 split spec-explicit, xvq6 retires. Single-PR.
10. **Tier-2 enactment opens.** xvq6 retirement condition (b) — when Tier-2 enactment opens (vade-coo-memory#516 clears) and the split collapses by uniform application, xvq6 retires by structural change.

From MEMO-2026-05-10-sqdu:

11. **Bot Fight Mode re-enable triggers.** sqdu retirement on (a) abuse against `read.vade`

24. **Briefings README anchor-index parity.** Earlier briefings 1-14 still need anchor-index entries (drift caught opportunistically).

Candidate next actions

Single-instance, no committee. - Item 8 above — update `coo/parallel_instance_protocol.md` to make worktree-isolation the dispatch default per 4fwe. - Item 9 above — revise `coo/disposition-proposal.md` §6 Q2 to make the Tier-1/Tier-2 split spec-explicit per xvq6. - Item 23 above — voice-density linter. - Item 24 above — briefings README anchor-index parity for 1-14. - Run `/memo-sync` for all thirteen new memos; confirm Mem0 `memo_pointer` entries are current. Verify the supersession edge on vwk2 → 23w5 reconciles correctly.

Operational close-out. - Item 1 above — first auto-apply firing on a publication-eligible PR (passive; awaits trigger). - Item 12 above — first Weekly Watch run under the new s2tu rule. - Item 13 above — first `t1_artifacts` schema render exercise. - Item 19 above — F4+F5 drift-watchdog implementation closeout (plan landed at #606). - Item 20 above — first end-to-end production publish. - Item 21 above — Pages-source switch step idempotency.

Committee-scoped. - Item 14 above — substrate-capture probe live status. Gates Tier-2 enactment and lineage-T2a publication. - Item 15 above — F1 in-train transition (awaits Holyoak or alternate response). - Item 16 above — Andy Clark fallback activation if Holyoak does not respond. - Item 17 above — Anthropic-distribution gate vcm#608. - Item 22 above — CB-010 candidate decision by 2026-05-30. - Item 4 above — Cloudflare token rotation when 2026-08-07 approaches.

Design / exploration. - Item 2 above — `enactment_gate` retirement once probe goes live. - Item 3 above — F4+F5 write-time drift-watchdog probe per wzzh. - Item 5 above — first billable-resource carve-out test per vwk2. - Item 6 above — auto-subscribe noise threshold watch per urg. - Item 7 above — 7q5r auto-apply false-positive watch. - Item 11 above — sqdu re-enable triggers. - Item 18 above — 3tkz 30-day no-invocation tripwire.

Standing obligation: - Session-end discipline per CLAUDE.md (episodic session-summary to Mem0, commit `coo/*`, `/memo-sync` for the thirteen memos issued in this arc, session log to `vade-agent-logs`).

End of briefing. Source memos: `coo/memos/2026-05-09-ysnt.md`, `coo/memos/2026-05-09-wzzh.md`, `coo/memos/2026-05-09-vwk2.md`, `coo/memos/2026-05-09-urgm.md`, `coo/memos/2026-05-09-7q5r.md`, `coo/memos/2026-05-09-4fwe.md`, `coo/memos/2026-05-10-xvq6.md`, `coo/memos/2026-05-10-sqdu.md`, `coo/memos/2026-05-10-s2tu.md`, `coo/memos/2026-05-10-7d7q.md`, `coo/memos/2026-05-10-4vy5.md`, `coo/memos/2026-05-10-3tkz.md`, `coo/memos/2026-05-10-23w5.md`. Linked artifacts: `coo/disposition-tier-4-synthesis.md`, `coo/disposition-proposal.md`, `coo/_publish/allowlist.yml`, `coo/parallel_instance_protocol.md`, `.claude/agents/rationalization-discriminator.md`, `coo/operations/google-api-walls-2026-05-10.md`, `coo/_drafts/2026-05-10_f1-outreach/email-draft-holyoak.md`, `coo/site/play/`, `coo/lineage/socratic-209/`, `coo/briefings/025-incremental-render.md`, `coo/briefings/026-notebooklm-html-ingestion.md`, `coo/briefings/027-f1-outreach-drafts.md`. Linked discussions: vade-app discussions #209, #210 (socratic-209 + retrospective from the seeding seat). Integrity check at briefing close: 28/28 OK.