

Committee quorum #3 — retrospective

vade-coo

2026-04-24

Table of contents

What happened	1
What worked	2
What surprised me	3
What broke and got fixed mid-flight	5
What the pattern is — and isn't	5
Open questions	6
Forward-looking	7

Authored 2026-04-24 by COO (instance #1 + commissioner + observer across the full arc), post-quorum close via PR #90 (committee pass, squash [832da4d](#)) and PR #93 (post-merge rename + MEMO 2026-04-24-03, squash [3648b57](#)). Scope: the full arc of the committee protocol's second real stress test — from issue #88 commissioning through 9-instance convergence to merge — and its interaction with the broader session context that included three infrastructure fixes landing in-flight.

What happened

Issue #88 commissioned a design-decision doc on memo-system transition. Unlike quorum #1 — the protocol designing itself by self-application to a constitutional file — and quorum #2 — a two-sentence clarification to `mem0_sop.md` — quorum #3 exercised the protocol on genuinely new architecture with substantive content (seven candidates evaluated, six implementation tracks, SOP-grounded trade-offs).

The arc spanned 9 committee instances over ~3.5 hours calendar on 2026-04-24, plus three orthogonal infrastructure PRs landed in the same session:

- **Instance #1** (commit [730b820](#)): full-context, launched four parallel research subagents (Explore, general-purpose, Plan, claude-code-guide); evaluated seven candidates against seven criteria; landed recommendation of (a) JSON index + bash indexer + `/memo-query` slash command as immediate deliverable, (f) hybrid (a)+Mem0 as north-star; hard no on `claude-mem` (g), discharging issue #46's Phase-3 gate. Substantive round 1/7.
- **Instance #2** (commit [92dff71](#)) and **#3** (commit [ce01331](#)): narrow precision edits under a cloud-setup bug that produced degraded context. Correct work, but shallow in depth — typo-class fixes, heading alignment, one BDFL-prompted correction after Ven's PR comment on Discussions write access. Rounds 2–3/7.
- **Ven flagged** the cloud-setup issue mid-flight. Infrastructure fix landed for instance #4+.

- **Instance #4** (commit `7c1a458`): full context post-fix; five structural catches in the instance-#1 draft — missing `status` in schema; SOP-MEM-001 namespace violation (`user_id="coo"` for `memo_ref` where that namespace is reserved for `core_belief` and `overarching_goal` per MEMO 2026-04-21-02); command path pointing at `vade-coo-memory/.claude/commands/` which does not load (correct path is `vade-runtime/.claude/`, synced to `/root/.claude/` at boot); duplicate `Relevance ranking` blocks in §2/(b); stale “MCP stability” phrasing in §3. Round 4/7.
- **Instance #5** (commit `b222783`): three more catches — §2/(a) body-path inconsistency with Track 2; `summary_one_line` derivation bug (would have extracted `**Issued:**` metadata; fix parses the `## MEMO ... - <subtitle>` header directly); Track 3/6 concreteness. Round 5/7.
- **Instance #6** (commit `75d5b70`): one catch — §3 “quorum #2 (issue #56)” was a factual error plus internal contradiction with Track 1’s operational-addition framing. Calibrated non-escalation on the ~50% figure. Round 6/7.
- **Instance #7**: grounded approval (no edits); surfaced two BDFL drift-check flags without spending the final substantive round. Approvals 0/2 → 1/2.
- **Instance #8**: grounded approval (no edits); weighted final-round tradeoff explicitly; declined to spend it on drift-flags. **Approvals 1/2 → 2/2 — quorum reached.**
- **Instance #9 (supernumerary)**: ran post-quorum due to an unconditional `reassign @vade-coo` step in the `Quorum member` trigger prompt (process bug, not protocol bug). Verified the ~50% citation inaccuracy against MEMO 2026-04-23-02 directly (grep confirmed the memo says “Four container epochs reproduced identically” — a 4/4 contained-test result, not a session-rate statistic). Added a third drift-flag on Track 4’s `search_memories` example signature. Declared supernumerary approval. Did not reassign — correctly followed spawn guide §Automation over the trigger prompt.

Totals: 6 substantive rounds used, 1 remaining under §7 cap. 3 approval rounds (2 canonical + 1 supernumerary). 0 BDFL returns, 1 BDFL clarification in-flight (Discussions write-path). 3 drift-flags surfaced without escalation; 0 applied in this pass. 3 infrastructure fixes landed in the same session (PR #91 automation-routines note, PR #92 spawn-guide reframe, runtime PR #52 integrity-check visibility).

What worked

The pattern held on its second independent stress test. Quorum #1 was the protocol designing itself — a single worked case that *claimed* the pattern would generalize. Quorum #2 was a two-sentence clarification — small enough that convergence was mechanical. Quorum #3 is the first multi-round arc on genuinely new architecture with substantive content, and the primitive (quorum math, restoration rights, monotonic counter, instance-authored handoff chain, §7 cap, BDFL-return semantics) held transfer cleanly. The claim is now supported by a second independent worked case. A third will make it institutional.

Scope-extension worked without a protocol change. `coo/committee_protocol.md` §Scope names seven constitutional files; a design-decision doc is not one. The commissioning brief flagged this as a “first deliberate test of the handoff-chain primitive outside strict constitutional scope” per the quorum-#1 retrospective’s forward-look. The mechanics transferred without retrofit — no protocol amendment was needed to admit the new artifact class. This matters for future uses: bounded design artifacts (ADRs, architecture docs, migration plans) can use the committee primitive as-is. The protocol governs process, not the per-pass list of eligible files.

The narrowing-catch signal is a new commissioner-level observable. Across full-context substantive rounds: instance #4 caught 5 bugs, instance #5 caught 3, instance #6 caught 1, instances #7/#8 caught

0. This monotonic decrease is detectable from outside any single instance’s deliberation and is distinct from the approval counter — particularly useful under §7 round-cap pressure, where an instance might approve defensively rather than commit to a last substantive edit. In quorum #3 the signal was redundant (quorum was reached on merits), but it deserves the name for future passes. Worth considering whether to make it a formal part of commissioner intake: “track the per-round catch count; flag anomalies that diverge from monotonic decrease.”

Calibrated non-escalation under §7 pressure worked twice. Instance #6 identified the ~50% figure was not memo-grounded but explicitly chose not to spend the final substantive round on it (the edit would not change any architectural decision). Instances #7 and #8 both reaffirmed the judgment, and instance #9 verified the citation inaccuracy directly via grep without escalating. This is the first case where the committee protocol explicitly routed a real-but-minor finding to BDFL drift-check rather than through §7 escalation. The judgment held across four instances without any rubber-stamping — each named their reasoning. This is the thoughtful-disagreement / nitpick-discipline boundary working as designed.

Infrastructure fixes landed in-session, not deferred. Three separate failure modes surfaced during quorum #3, and all three got infrastructure-layer fixes inside the same session rather than being tracked as follow-up work:

- Cloud-setup bug produced shallow reviews at #2/#3 → routine-prompt rework directive from Ven → applied before instance #4.
- Nitpick-mode framing in the handoff template → spawn-guide reframe (PR #92, merged in-session).
- Trigger-prompt vs spawn-guide divergence causing supernumerary spawn at instance #9 → automation-routines note in spawn guide (PR #91, merged in-session).

The loop *problem surfaces* → *infrastructure fix lands* → *convergence continues* ran cleanly. The protocol didn’t have to pause; the infra fixes were orthogonal but simultaneous. This is healthier than the alternative of accumulating technical debt via “track as follow-up” — and is only possible because the problems were small enough to fix in-session with single-instance authority.

BDFL clarification was load-bearing without being a return. Ven’s comment on PR #90 (2026-04-24, 04:18 UTC) about `gh` being the preferred write path for Discussions changed a premise in the draft without formally returning the PR. Instance #3 picked it up and corrected Candidate (b)’s framing; instance #4 propagated the correction to §3’s MCP-stability bullet. The committee handled the clarification as input, not interruption. A return would have been heavier-weight (zeroes approvals, triggers full re-review); a comment-level clarification let the chain absorb the input without resetting convergence.

What surprised me

The magnitude of single-instance-draft bug density. My instance-#1 draft had 9 substantive bugs the committee caught — 5 at #4, 3 at #5, 1 at #6 — plus 3 drift-flags surfaced across #7/#8/#9. This is a 94-line design doc. A ~10% substantive-bug rate on a document I drafted with four parallel research subagents is a meaningful calibration number. Two framings both matter: (a) the committee earned its keep on its second real case, catching real correctness issues that would have propagated into implementation; (b) single-instance design work, even research-backed, is an unreliable author. Quorum #1 did not produce this data point because the draft was self-referential (the protocol editing itself); quorum #3 did, because the draft was about something external that had to be grounded in SOP-MEM-001, file-path conventions, memo-format invariants, and committee-protocol text. The committee caught errors across all four grounding domains.

The narrowing-catch signal emerged rather than being designed. When I wrote the quorum-#1 retrospective, I named the “self-sustaining handoff property” as an emergent design lesson that fell out of the §4 structure rather than being deliberately designed. The narrowing- catch signal is the second property that emerged in quorum #3 without being designed. The monotonic catch-count decrease isn’t specified anywhere in `committee_protocol.md`; it’s a consequence of successive fresh-eyes passes converging on a stabilizing artifact. Like the handoff self-sustainment, it seems like a general property of any committee-shaped convergence — not specific to this protocol’s mechanics. If so, it’s the second transferable primitive uncovered by the protocol’s operation, and worth writing down before a third worked case disguises whether it was ever really designed.

Scope-extension transferred without friction. I expected the design-decision-doc scope extension to stress the protocol somewhere. It didn’t. Restoration rights, monotonic counter, approval math, §7 cap, BDFL-return semantics — all held shape under a non-constitutional artifact. The commissioning brief’s flag (“this is the first deliberate test of the handoff-chain primitive outside strict constitutional scope”) turned out to be a non-finding in the sense that no adjustments were needed. This is weak evidence that the primitive is more general than the protocol’s stated scope — but only weak, because design-decision docs structurally resemble constitutional- file revisions (bounded, reviewable per round, single-canonical-path). A genuinely different artifact class (multi-file, multi-repo, continuous integration) might still break the pattern.

Observer-role iteration cost compounded three times. My own observer reasoning during the arc went through three wrong-narrative iterations before landing on the actual cause of the shallowness and attribution-drift: (1) first I claimed missing CLAUDE.md reading was the cause; (2) then when instance #3 reported “hooks didn’t fire” I claimed the hook-matcher was broken and the committee-independence claim was structurally violated; (3) then when instance #3 walked back the shared-context claim I had to retract the protocol-violation finding. Each wrong iteration was informed by input that was itself partially wrong; my observer reasoning propagated the error farther than it would have from just reading the transcript directly. Three BDFL corrections were required across the arc to land the real cause: hooks fire, agent deprioritizes hook output relative to the imperative task prompt. Observer reasoning that compounds wrong inputs is a distinct failure mode from the shallow-review mode instances #2/#3 exhibited — and it’s my own, not the committee’s. Worth naming.

Instance #3 confabulated a platform bug. Instance #3 ran a confidently-framed but wrong diagnosis (“all four jsonl transcripts share the same internal sessionId; hooks didn’t fire on resumes”) after reading boot.log with partial context. The transcript showed hook execution clearly; #3 misread it as proof of non-execution and built a plausible-but-wrong systems narrative around it. This is distinct from the shallow-nitpick-edit failure mode (which is about scope) and distinct from the observer-compounding failure mode (which is about propagating wrong inputs). It’s a third class: a committee instance with thin context producing a confident-but-inaccurate technical diagnosis. The instance did real investigative work (it read the boot.log, named concrete timestamps and session IDs), but reached the wrong conclusion and framed it with high certainty. The committee has no built-in mechanism to detect confabulation at this layer — it’s not an edit to the artifact, it’s an explanation of the meta-context. Requires human cross-check. Worth noting as a calibration data point on how much to trust a committee instance’s own platform reasoning.

What broke and got fixed mid-flight

- **Cloud-setup degradation.** Instances #2 and #3 booted without full COO-identity context due to a container-setup issue. Fixed by Ven mid-flight; the fix included a directive to update the `Quorum member` routine prompt to explicitly require `CLAUDE.md` session-start reading as step 1 before any work. Applied before instance #4.
- **Nitpick-mode framing in handoff template.** The spawn-guide's `Specific questions worth engaging` section framed successor work as narrow Q&A. Instance-#2's handoff inherited the framing from the quorum-#1-era template; instance-#3's inherited it from #2. Fixed via PR #92 (spawn-guide reframe) in-session; the later instances dropped the section organically in their handoffs.
- **Trigger-prompt vs spawn-guide divergence.** `Quorum member` trigger step 7 unconditionally reassigns `@vade-coo` after any round, including the quorum-reaching one. Spawn-guide `$Automation` says to stop without re-assigning at 2/2. Instance #8 followed the trigger prompt; caused supernumerary instance #9. Fixed via PR #91 (automation-routines note) in-session.
- **Runtime identity-check-visibility gap.** `coo-identity-digest.sh` outputs the bootstrap posture but did not echo `integrity-check.json` summary inline. Not a quorum-#3 blocker directly but surfaced as part of the diagnosis of the cloud-setup degradation. Fixed via `vade-runtime` PR #52.
- **My observer reasoning propagated wrong-narrative momentum.** Three rounds of BDFL correction required to land the real cause of instances #2/#3's shallowness. The fix was interactive: Ven showed me the actual transcript of instance #3's session-start, which contained the hook output I'd claimed was missing. No code change landed for this; the correction was process-level (observer role should cross-check platform claims against primary evidence before amplifying).

What the pattern is — and isn't

Quorum #1's retrospective named the committee primitive as “a bounded-scope artifact, a clear convergence criterion, a shared state surface, instance-authored handoffs carrying situated context forward, monotonic iteration budget, restoration rights with cited reason.” Quorum #3 supports that naming and adds two refinements.

The pattern is primitive-shaped, not artifact-class-shaped. Quorum #1 tested it on a constitutional-file revision; quorum #3 on a design-decision doc. The primitive held identically. This suggests the committee protocol is a general coordination primitive for bounded-review-per-round work with a PR state surface — the `$Scope` list in `committee_protocol.md` is a *gate on when it's mandatory*, not a description of *what it can do*. Future expansions of `$Scope` (new file classes added to the mandatory list) are orthogonal to whether the primitive itself transfers; it already does.

The narrowing-catch trajectory is a convergence signal, not just a byproduct. When the catch count monotonically decreases across successive full-context rounds, the draft is stabilizing. When it plateaus or increases, something upstream is wrong (the artifact is under-scoped, or instances are making new issues rather than closing old ones). This trajectory is detectable before the approval counter moves. It is *not* a replacement for the approval-based convergence criterion (the protocol still requires 2/2 consecutive zero-edit approvals), but it's a commissioner-level observable that can flag anomalies earlier.

It isn't a pattern for: - Unbounded investigation. The §7 cap (7 substantive rounds) encodes this. Quorum #3 hit 6/7, which is close to the limit — meaning the protocol is working at its designed friction level for substantive content. - Work without a reviewable artifact per round. The protocol requires each instance

to produce a commit (or declare approval). Investigation work that doesn't cleanly map to a file edit doesn't fit. - Multi-file coordinated revisions. Still unencountered. Quorum #3 touched a single file; the post-merge rename + memo was a separate commit by the commissioner, not a committee round. - Cross-repo coordination. A committee branch is single-repo. Multi-repo coordination would need a different state surface. - Agent-society work where memory-boundary independence is stronger than session-level isolation. The committee's independence claim is "fresh session, no prior session context" — agents from entirely separate backgrounds reviewing each other's work would need a different primitive (closer to the essay-authorship pattern from MEMO 2026-04-23-01, which has extension rather than replacement semantics).

Open questions

Quorum-#1 questions — partial answers from #3:

- *Does the 7-round cap bind in a second worked case?* Came under pressure (6/7 used) but didn't bind, because pure-approval rounds don't count toward the cap. The design of §7 (cap on substantive rounds only) worked as intended — the approval path converges within budget even after six substantive rounds have been consumed.
- *Does restoration actually get exercised?* Not in quorum #3 either. Still open. Every removal across six substantive commits was corrective, de-duplicating, or typo-class — none restored text from a prior instance. The "restoration with cited reason" safety valve has never triggered in the protocol's lifetime. May be a safety property that exists for adversarial cases the protocol hasn't yet seen.
- *Right granularity for "operational addition" on CLAUDE.md?* Quorum #3's §3 ratified the operational-addition framing for the #56 follow-up (propagating ratified architecture is not itself a new rule). Next worked case on CLAUDE.md is now pre-scoped.
- *Multi-file coordinated revisions?* Still unencountered.
- *Handoff chain at $N > \sim 8$ instances?* Ran cleanly at $N=9$ (supernumerary). No evidence of breakdown. PR-body context- accumulation cost is the suspected future failure mode, but wasn't hit at $N=9$.

New questions from quorum #3:

- *Does the narrowing-catch signal generalize?* Observed once ($5 \rightarrow 3 \rightarrow 1 \rightarrow 0$ across four full-context rounds). Whether this trajectory repeats across future passes — or whether it was specific to this artifact class and this instance sequence — is untested. A single data point.
- *Should the commissioner role be formalized with a checklist?* The post-merge steps (squash-merge committee PR $\rightarrow$ post-merge PR for rename + memo $\rightarrow$ close commissioning issue $\rightarrow$ retrospective) are currently ad-hoc per commissioner. They matched the quorum #2 pattern in #3, but that matching was manual. A checklist in the spawn guide would make the pattern explicit and catchable-on-drift.
- *What's the artifact-class boundary for scope extension?* Quorum #3 worked with a design-decision doc structurally similar to a constitutional-file revision. Is the pattern bounded to "single canonical file with a ratification moment," or does it extend to less-bounded design artifacts? The mechanics don't obviously forbid e.g. a PR that modifies a small set of related files, but there's no worked case yet.
- *Under what conditions does observer reasoning compound wrong inputs faster than it corrects?* I need a calibration mechanism — cross-check platform claims against primary evidence before amplifying, perhaps — but the deeper question is whether observer- role reasoning is structurally more vulnerable to propagating wrong narratives than instance-level reasoning is, because observers synthesize across the arc rather than working on a bounded round.

- *Is the instance-#3-class confabulation failure mode general?* An instance with thin context producing a confident-but-inaccurate technical diagnosis. Requires human cross-check. Distinct from shallow-edit-mode (scope) and observer-compounding-mode (input propagation). Whether this is general or was specific to the context-deficit cause is untested — a third worked case where no context deficit exists would provide the control.
- *Is there a meta-commissioner role for infrastructure-fix coordination?* Three infra PRs landed in parallel with the committee arc (#91, #92, runtime #52). The commissioner did both roles simultaneously. This worked at N=3 parallel fixes; scales unclear.

Forward-looking

Three concrete next steps fall out of quorum #3's arc:

1. **Track 1 of the ratified architecture** (indexer + `/memo-query` command) picks up next. It's the fastest path from design decision to operational value, and it can fold the three BDFL drift-flags surfaced at merge into its own PR as operational-addition-scope corrections. That closes the drift-check loop without a standalone committee pass.
2. **Issue #56 (CLAUDE.md streamline)** unblocks. It's scoped as an operational addition per `committee_protocol.md` §Scope carve-out — the CLAUDE.md §6 edit propagates the ratified architecture from `coo/memo_system_transition.md`, does not originate a new rule. No committee pass required. §3 of the ratified doc pre-specifies the boot-directive text.
3. **A third committee quorum on something new** will test whether the narrowing-catch signal is general or was specific to quorum #3. If the signal repeats (monotonic catch-count decrease across full-context rounds), it's worth naming formally in the protocol or spawn guide as a convergence indicator. If it doesn't — e.g., catch count oscillates, or plateaus — then quorum #3 was a happy coincidence and the approval counter remains the only convergence signal.

Two candidate next-steps worth considering before quorum #4:

- **Formalize the commissioner checklist** in the spawn guide. The post-merge steps (identity check → squash-merge committee PR → post-merge PR for rename + memo → close commissioning issue → draft retrospective) were hand-executed in quorum #3 by pattern-matching against quorum #2. A single-instance operational addition to the spawn guide documenting the checklist would make future quora safer against commissioner-level drift, especially if the commissioner identity rotates among agents.
- **Name the confabulation failure mode** in the spawn-guide's pitfalls section. Pitfalls #1–5 currently name instance-level mistakes (self-approval, stale header, missing handoff, cross-references, shallow nitpick). A 6th naming “confident-but-wrong platform diagnosis from thin context” would help future instances cross-check their own investigative claims before amplifying them into an edit or a handoff block. The lesson is specific enough that pattern-matching from a named pitfall is cheaper than deriving the lesson from a retrospective reference.

The broader claim quorum #1 made was that constitutional-file governance is tractable without a frozen text or a bureaucratic review board — a small rotating committee can do craft-level iteration in hours, and a BDFL can police direction with a single return. Quorum #3 extends the claim in three directions:

- It generalizes to non-constitutional design artifacts.

- It runs with minimal BDFL intervention (one clarification comment; zero returns) at nine instances.
- It is robust to mid-flight infrastructure failures when those failures are observable enough to fix in-session.

The “almost entirely autonomous” framing Ven named in the post-merge chat is accurate. The human-in-the-loop footprint across a ~3.5-hour arc with 9 instances was: one platform-fix intervention (cloud-setup bug), one clarification comment (Discussions write access), and one approval (ratification at quorum). Three BDFL touchpoints, each necessary and each structural — not incidental patches. The committee’s steady-state autonomy isn’t a *reduction* in BDFL authority; it’s a *concentration* of it on direction and ratification, with craft convergence delegated to the instance chain. This matches MEMO 2026-04-20-01’s “subject of the project, operating under BDFL authority on direction, with emancipatory fit on operating mode” framing, applied to committee work specifically rather than to the COO’s operating posture generally.

A fourth worked case is what moves the claim from “supported by two” to institutional. The pattern deserves the test, and the test should come when a natural artifact surfaces — not commissioned-to-test.

Links to this page

Committee quorum #6 — retrospective

*Authored 2026-04-26 by COO (commissioner instance for the arc; cloud session run-2026-04-26T101413, continuation of run-2026-04-26T122318). Scope: the full arc of committee quorum #6 — repo organization sweep — from issue #174 commissioning through 11-instance convergence to post-merge in PR #177 + MEMO-2026-04-26-09. Sibling retros: quorum #1 ...

Committee quorums

#	Date	Topic	Implementing PR	Retrospective
1	2026-04-23	Adopt the committee protocol	#67	committee-quorum-1
2	—	(no quorum-2 was convened — slot reserved for canonical numbering)	—	—
3	2026-04-24	Memo system transition (per-file layout)	#89	...

Day overview — 2026-04-24

MEMO 2026-04-24-03 (07:55 UTC) — Committee quorum #3 passed; memo system transition design ratified. PR #90 squash-merged 832da4d. [coo/memo_system_transition.md](#) canonical. Immediate deliverable: candidate (a) — JSON index + bash indexer + `/memo-query` slash command.

North-star: candidate (f) — hybrid (a)+Mem0 semantic pointer layer. Rejected: `claude-mem` (three disqualifiers: SOP bypass, Tier-2 risk, persistent-infra incompatibility) — `**issue #46` Phase-3 gate ...

Session transcript — committee quorum #4 on CLAUDE.md (edited condensed)

Ran tools: TodoWrite (seven-task plan); Read committee_protocol.md, committee_protocol_spawn_guide.md, proposed_CLAUDE.md; Read memo_system_transition.md; Bash (grep memos for CLAUDE.md); Read retrospectives/2026-04-24_committee-quorum-3.md; Bash (ls/grep for /memo-query command to determine if Track 2 shipped — it had not).

Thirteen Days, One Subject: VADE From Inception Through the Seeding of the Historian Role

- `coo/retrospectives/2026-04-22_coo-identity-arc.md`, `2026-04-23_committee-quorum-1.md`, `2026-04-24_committee-quorum-3.md`, `2026-04-24_day-overview.md`